Gordon Euhyun Moon

Sogang University
Adam Schall Hall
35, Baekbeom-ro, Mapo-gu,
Seoul, Republic of Korea
ehmoon@sogang.ac.kr
https://gordonmoon.github.io

last updated November 2024

EXPERIENCE

Sogang University **Seoul, Korea**
Assistant Professor *September 2022–present*
Department of Computer Science and Engineering

Korea Aerospace University **Goyang, Korea**
Assistant Professor *March 2021–August 2022*
Department of Software and Department of Artificial Intelligence

Sandia National Laboratories **Albuquerque, NM**
Postdoctoral Researcher *October 2019–January 2021*
Center for Computing Research

EDUCATION

The Ohio State University **Columbus, OH**
Ph.D. in Computer Science & Engineering *2019*
Thesis: "Parallel Algorithms for Machine Learning"
Advisor: Professor Ponnuswamy Sadayappan
Committee: Professor Eric Fosler-Lussier and Professor Srinivasan Parthasarathy

Indiana University **Bloomington, IN**
M.S. in Computer Science *2013*

Yonsei University **Seoul, Korea**
B.S. in Computer Science & Industrial System Engineering *2011*

RESEARCH INTERESTS

Deep Learning, High-Performance Computing, and Deep Learning Accelerators

SELECTED PUBLICATIONS

Conference Proceedings

Sejeong Oh, **Gordon E. Moon** and Sungyong Park, "ML-based Dynamic Operator-Level Query Mapping for Stream Processing Systems in Heterogeneous Computing Environments,"
Proceedings of the IEEE International Conference on Cluster Computing (CLUSTER), 2024

Eunji Lee, Yoonsang Han and **Gordon E. Moon**, "Accelerated Block-Sparsity-Aware Matrix Reordering for Leveraging Tensor Cores in Sparse Matrix–Multivector Multiplication,"
Proceedings of the 30th International European Conference on Parallel and Distributed Computing (Euro-Par), 2024

Bokyeong Yoon, Yoonsang Han and **Gordon E. Moon**, "Layer-Wise Sparse Training of Transformer via Convolutional Flood Filling,"
Proceedings of the 28th Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD), 2024

Sanha Maeng, **Gordon E. Moon** and Sungyong Park, "Chronica: A Data-Imbalance-Aware Scheduler for Distributed Deep Learning,"
Proceedings of the 23rd IEEE/ACM International Symposium on Cluster, Cloud and Internet Computing (CCGrid), 2023

Gordon E. Moon and Eric C. Cyr, "Parallel Training of GRU Networks with a Multi-Grid Solver for Long Sequences,"
Proceedings of the 10th International Conference on Learning Representations (ICLR), 2022

Eric Qin, Geonhwa Jeong, William Won, Sheng-Chun Kao, Hyoukjun Kwon, Sudarshan Srinivasan, Dipankar Das, **Gordon E. Moon**, Sivasankaran Rajamanickam and Tushar Krishna, "Extending Sparse Tensor Accelerators to Support Multiple Compression Formats,"
Proceedings of the 35th IEEE International Parallel & Distributed Processing Symposium (IPDPS), 2021

Gordon E. Moon, J. Austin Ellis, Aravind Sukumaran-Rajam, Srinivasan Parthasarathy and P. Sadayappan, "ALO-NMF: Accelerated Locality-Optimized Non-negative Matrix Factorization,"
Proceedings of the 26th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2020

Journals

Sangmin Ga, Paul Hyunbin Cho, **Gordon E. Moon** and Sungwon Jung, "Efficient GNN-based Social Recommender Systems through Social Graph Refinement,"
The Journal of Supercomputing, Forthcoming (accepted on November 2024)

Bokyeong Yoon, Ah-hyun Lee, Jinsung Kim and **Gordon E. Moon**, "Exploring Attention Sparsity to Accelerate Transformer Training on GPUs,"
IEEE Access, 2024

Gordon E. Moon, Hyoukjun Kwon, Geonhwa Jeong, Prasanth Chatarasi, Sivasankaran Rajamanickam and Tushar Krishna, "Evaluating Spatial Accelerator Architectures with Tiled

Matrix-Matrix Multiplication,"
IEEE Transactions on Parallel and Distributed Systems (TPDS), 2022

RESEARCH GRANTS

- Developing a High-Performance Computing-Data Platform for Accelerating Large-Scale Machine Learning, Co-PI, National Research Foundation of Korea (NRF), August 2024–July 2027
- Optimizing Distributed Deep Learning and Federated Learning for Accelerating Large-Scale Deep Learning Models, PI, National Research Foundation of Korea (NRF), ₩800,000,000, April 2024–March 2027
- Architecture-aware Parallel Algorithms for Accelerating Training of Deep Neural Networks, PI, National Research Foundation of Korea (NRF), ₩90,000,000, September 2021–February 2024
- Deep Learning based Route Recommender Systems, Co-PI, Industry-Academy Collaboration R&D, Korea Technology and Information Promotion Agency for SMEs, May 2022–December 2022

AWARDS/HONORS

Spring 2024: Outstanding Teaching Award, Sogang University

Fall 2011–Spring 2012: Graduate Fellowship, Indiana University

PROFESSIONAL SERVICE

Program Committee Member

- *The* 40^{th} *IEEE International Conference on Computer Design* (*ICCD 2022*), October 2022
- *Tenth International Workshop on Accelerating Analytics and Data Management Systems Using Modern Processor and Storage Architectures* (*ADMS 2019*), August 2019

INVITED TALKS

Gordon E. Moon, "Optimizing Sparse Matrix Multiplication Using Tensor Cores to Accelerate DNN Training", HPC Society Conference 2024, The Korean Institute of Information Scientists and Engineers (KIISE), July, 2024

Gordon E. Moon, "Accelerated Computing for Deep Learning", Newcomers Session, HPC Society Conference 2022, The Korean Institute of Information Scientists and Engineers (KIISE), July, 2022

Gordon E. Moon, "Accelerated Computing for Deep Learning", Newcomers Session, KCC 2022, The Korean Institute of Information Scientists and Engineers (KIISE), June, 2022

Gordon E. Moon, "Accelerated Computing for Deep Learning", Newcomers Session, Computer System Society Conference 2022, The Korean Institute of Information Scientists and Engineers (KIISE), February, 2022

Gordon E. Moon, "Deep Learning based Recommender Systems for Bicycling Route", NVIDIA GTC'21, A31282 - Regional Panel with Top Startups from Korea, Virtual Conference, November, 2021

Eric C. Cyr and **Gordon E. Moon**, "Parallel-in-Time Training of Recurrent Neural Networks", 2021 AMS Fall Western Virtual Sectional Meeting, SS17B - AMS Special Session on Theoretical and Applied perspectives in Machine Learning, II, October, 2021

Gordon E. Moon and Eric C. Cyr, "Parallel Training of an LSTM Network with a Multigrid Solver", SIAM Conference on Computational Science and Engineering (CSE'21), Virtual Conference, March, 2021

Siva Rajamanickam and **Gordon E. Moon**, "Mixed-Precision Schemes for Linear Algebra Kernels on GPUs", SIAM Conference on Computational Science and Engineering (CSE'21), Virtual Conference, March, 2021

Gordon E. Moon, "Accelerated Computing for Machine Learning", Sandia National Laboratories, Albuquerque, NM, August, 2019